

What makes a model?

Fitting Models

How do you fit a linear model in R?

How many different ways can you think of?

- `lm` for linear model
- `glm` for generalized linear model (e.g. logistic regression)
- `glmnet` for regularized regression
- `keras` for regression using TensorFlow
- `stan` for Bayesian regression
- `spark` for large data sets

To specify a model

- Choose a model
- Specify an engine
- Set the mode

To specify a model

```
1 library(tidymodels)
2 linear_reg()
3 #> Linear Regression Model Specification (regression)
4 #>
5 #> Computational engine: lm
```

To specify a model

```
1 linear_reg() %>%  
2   set_engine("glmnet")  
3 #> Linear Regression Model Specification (regression)  
4 #>  
5 #> Computational engine: glmnet
```

To specify a model

```
1 linear_reg() %>%  
2   set_engine("stan")  
3 #> Linear Regression Model Specification (regression)  
4 #>  
5 #> Computational engine: stan
```

To specify a model


```
1 decision_tree()  
2 #> Decision Tree Model Specification (unknown mode)  
3 #>  
4 #> Computational engine: rpart
```

To specify a model


```
1 decision_tree() %>%  
2   set_mode("regression")  
3 #> Decision Tree Model Specification (regression)  
4 #>  
5 #> Computational engine: rpart
```

All available models are listed at <https://www.tidymodels.org/find/parsnip/>


Linear regression



Linear regression



Linear regression




- Outcome modeled as linear combination of predictors:

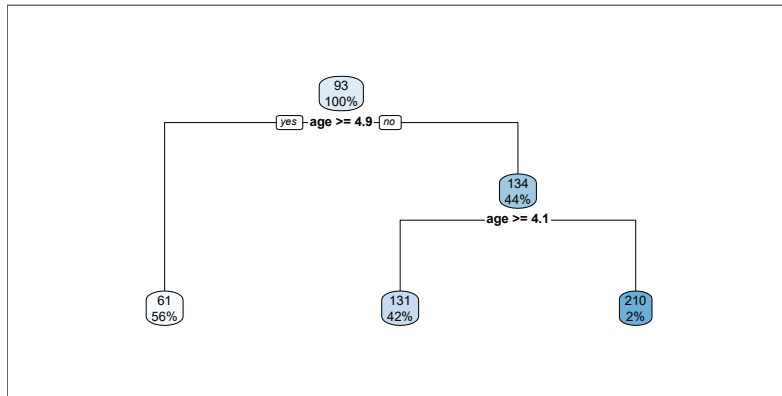
$$\text{latency} = \beta_0 + \beta_1 \cdot \text{age} + \epsilon$$

- Find a line that minimizes the mean squared error (MSE)

Decision trees


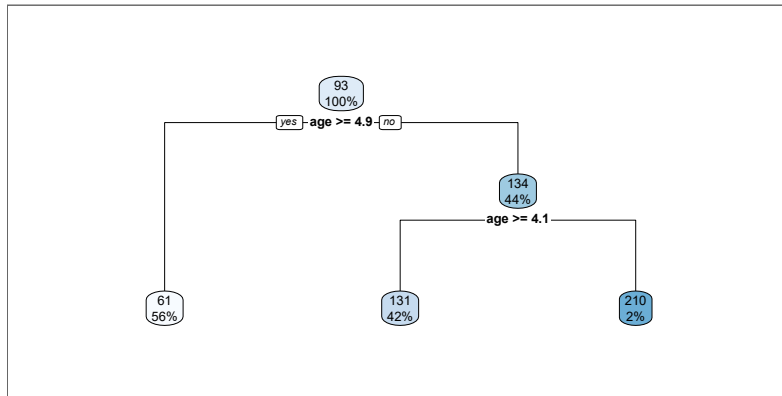


Decision trees




- Series of splits or if/then statements based on predictors
- First the tree *grows* until some condition is met (maximum depth, no more data)
- Then the tree is *pruned* to reduce its complexity

Decision trees




All models are wrong, but some are useful!

Linear regression

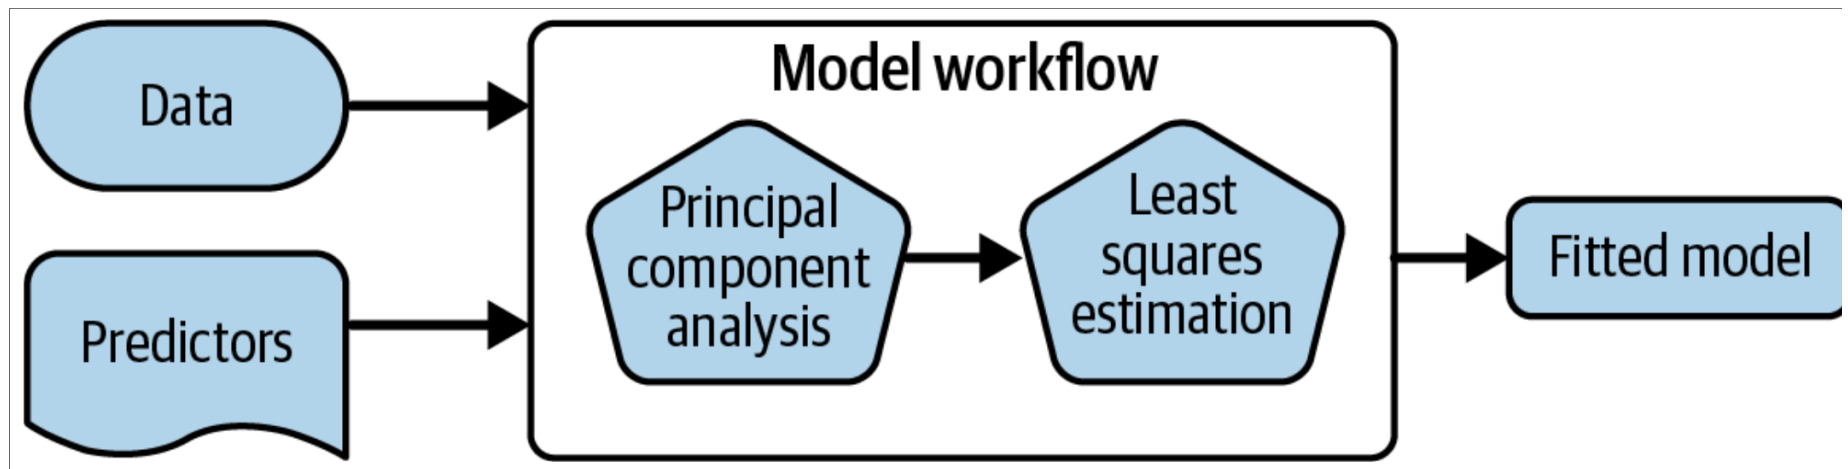


Decision trees

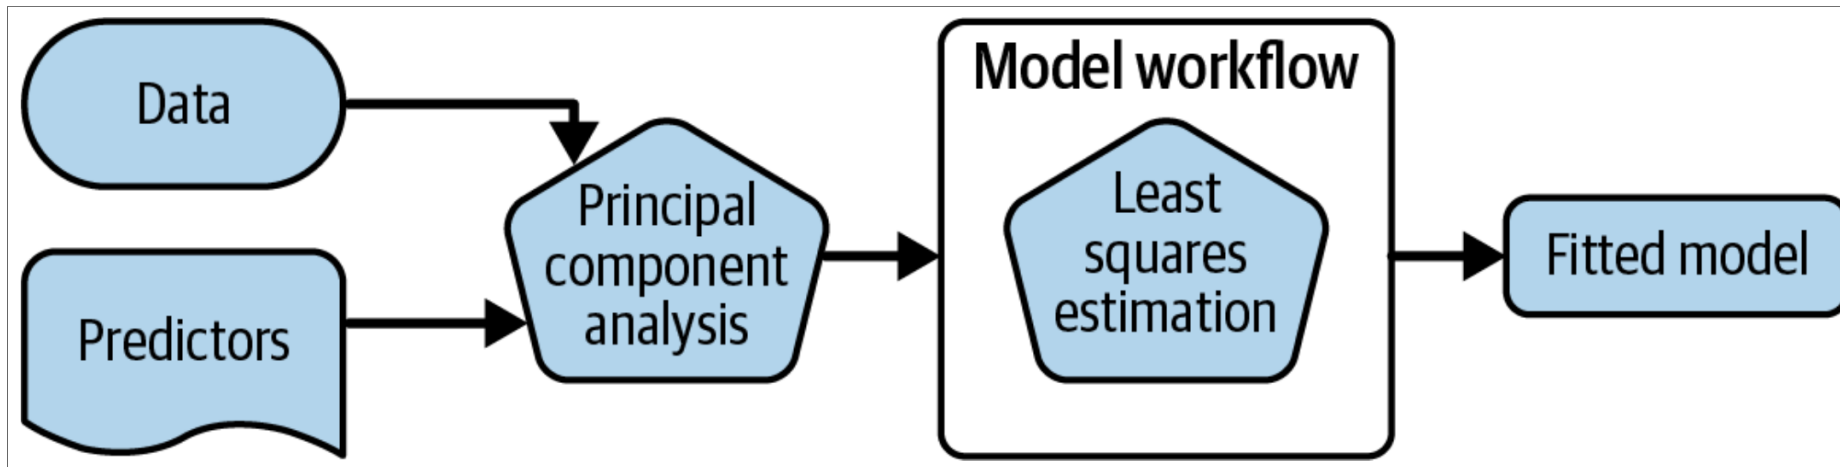


A model workflow

Workflows bind preprocessors and models



What is wrong with this?



Why a `workflow()`?

- Workflows handle new data better than base R tools in terms of new factor levels
- You can use other preprocessors besides formulas (more on feature engineering tomorrow!)
- They can help organize your work when working with multiple models
- Most importantly, a workflow captures the entire modeling process: `fit()` and `predict()` apply to the preprocessing steps in addition to the actual model fit

A model workflow

```
1 tree_spec <-  
2   decision_tree() %>%  
3   set_mode("regression")  
4  
5 tree_spec %>%  
6   fit(latency ~ ., data = frog_train)  
7 #> parsnip model object  
8 #>  
9 #> n= 456  
10 #>  
11 #> node), split, n, deviance, yval  
12 #>      * denotes terminal node  
13 #>  
14 #> 1) root 456 2197966.00  92.90351  
15 #> 2) age>=4.947975 256  252347.40  60.89844  
16 #> 4) treatment=control 131  91424.06  48.42748 *  
17 #> 5) treatment=gentamicin 125  119197.90  73.96800 *  
18 #> 3) age< 4.947975 200 1347741.00 133.87000  
19 #> 6) treatment=control 140  986790.70 118.25710
```

A model workflow

```
1 tree_spec <-  
2   decision_tree() %>%  
3   set_mode("regression")  
4  
5 workflow() %>%  
6   add_formula(latency ~ .) %>%  
7   add_model(tree_spec) %>%  
8   fit(data = frog_train)  
9 #> == Workflow [trained] ==  
10 #> Preprocessor: Formula  
11 #> Model: decision_tree()  
12 #>  
13 #> — Preprocessor —  
14 #> latency ~ .  
15 #>  
16 #> — Model —  
17 #> n= 456  
18 #>  
19 #> node), split, n, deviance, yval
```

A model workflow

```
1 tree_spec <-  
2   decision_tree() %>%  
3   set_mode("regression")  
4  
5 workflow(latency ~ ., tree_spec) %>%  
6   fit(data = frog_train)  
7 #> == Workflow [trained] ==  
8 #> Preprocessor: Formula  
9 #> Model: decision_tree()  
10 #>  
11 #> — Preprocessor —  
12 #> latency ~ .  
13 #>  
14 #> — Model —  
15 #> n= 456  
16 #>  
17 #> node), split, n, deviance, yval  
18 #>      * denotes terminal node  
19 #>
```

Predict with your model

How do you use your new `tree_fit` model?

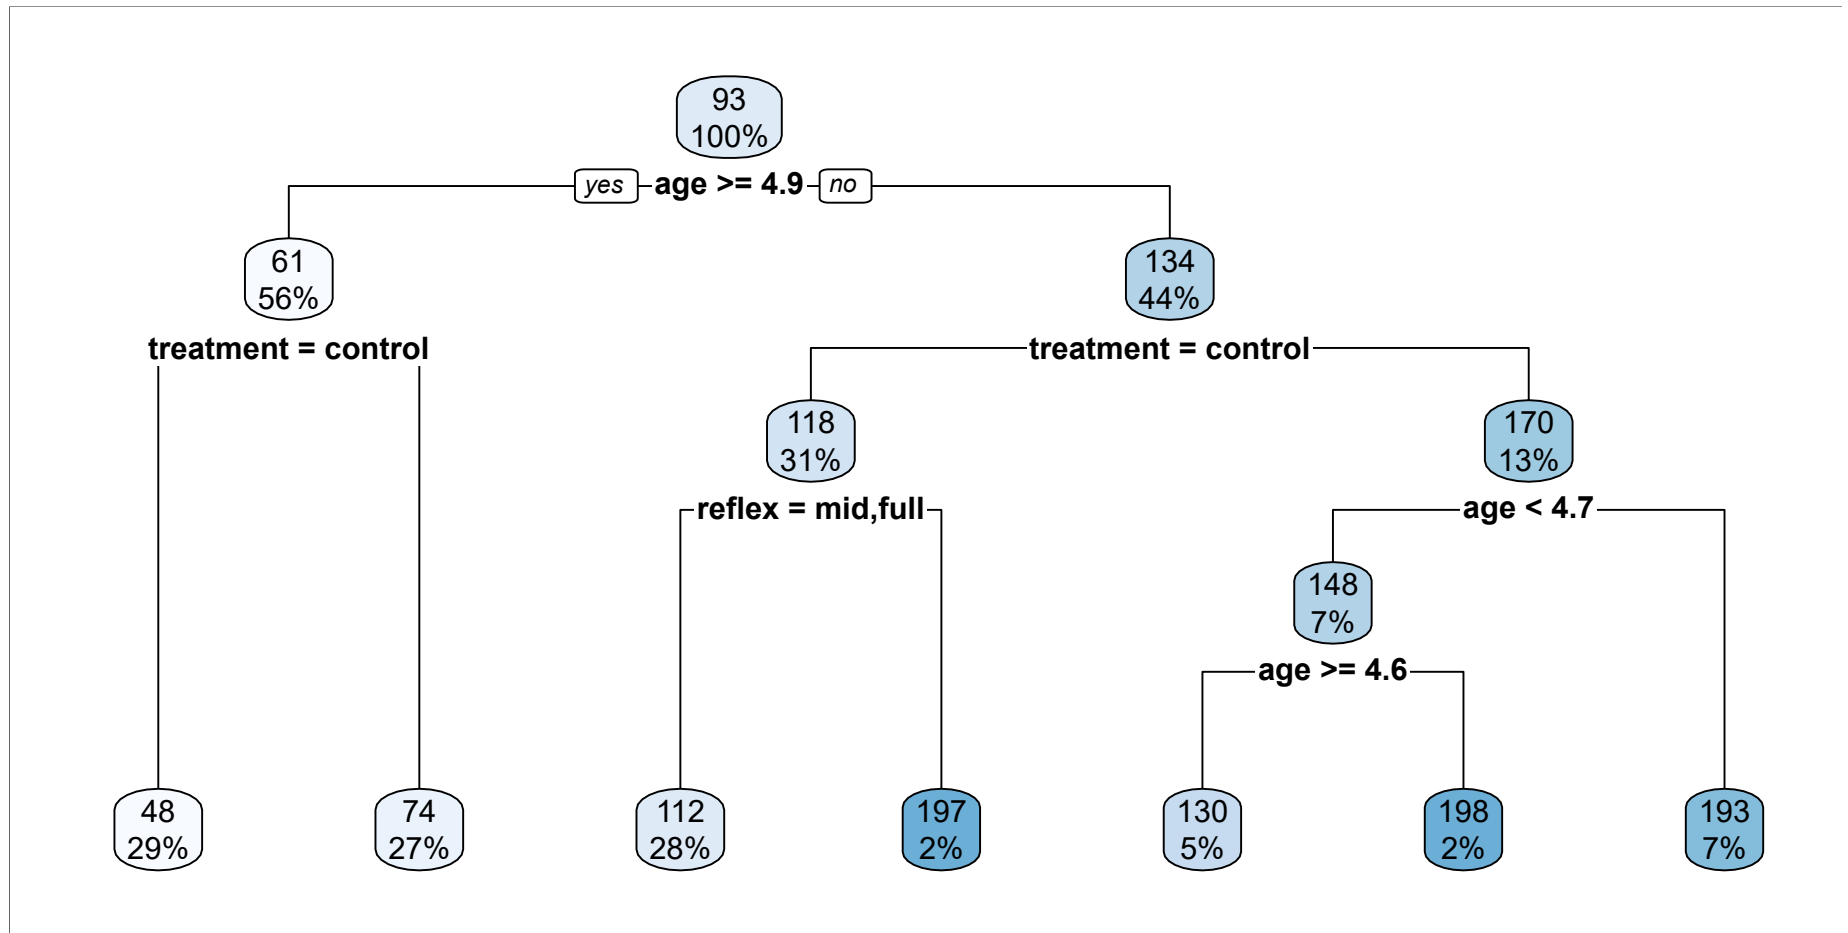
```
1 tree_spec <-  
2   decision_tree() %>%  
3   set_mode("regression")  
4  
5 tree_fit <-  
6   workflow(latency ~ ., tree_spec) %>%  
7   fit(data = frog_train)
```


The tidymodels prediction guarantee!

- The predictions will always be inside a **tibble**
- The column names and types are **unsurprising** and **predictable**
- The number of rows in **new_data** and the output **are the same**

Understand your model

How do you understand your new `tree_fit` model?



Understand your model

How do you understand your new `tree_fit` model?

```
1 library(rpart.plot)
2 tree_fit %>%
3   extract_fit_engine() %>%
4   rpart.plot(roundint = FALSE)
```

You can `extract_*()` several components of your fitted workflow.

Understand your model

How do you **understand** your new `tree_fit` model?

You can use your fitted workflow for model and/or prediction explanations:

- overall variable importance, such as with the `vip` package
- flexible model explainers, such as with the `DALEXtra` package

Learn more at <https://www.tmwr.org/explain.html>

Error

[×](#)